

《强化学习》专业选修课

作业一

提交时间：10 月 15 日之前

提交形式：请提交一个压缩包，命名为：学号_姓名_assignment1.zip。内容包括：

- 题目 1 和 2 的解答 assignment1_solution.pdf
- 题目 3 中的 algorithm1.py
- Email 到：bit_reinforce@163.com，邮件标题为：学号_姓名_assignment1

题目 1：奖励函数的选择

已知一个离散动作状态空间 MDP，折扣因子 $\gamma \in (0,1)$ 。该 MDP 没有终止状态，智能体将不断做出决策。假设该问题的原始最优值函数为 V_1^* ，最优策略为 π_1^* 。

- 若对该 MDP 的每一次状态转移上都加上一个小的正实数（即： $r(s,a)$ 是原始奖励函数，新的奖励函数为 $\hat{r}(s,a) = r(s,a) + c; \forall s,a$ ），请写出新的最优值函数表达式。在新条件下，最优策略是否会改变？请说明理由。
- 若对该 MDP 的每一次状态转移都乘上实数 $c \in R$ （即 $\hat{r}(s,a) = c \times r(s,a); \forall s,a$ ）。
 - 是否存在特定情况使得新的最优策略仍然为 π_1^* ，且值函数能用 c 和 V_1^* 来表示？如果存在，请写出对应的值函数表达式、满足该情况的常数 c 。如果不存在，说明理由。
 - 是否存在特定情况使得最优策略发生变化？如果存在，请举出一个 c 的例子，并说明变化理由；如果不存在，也请说明理由。
 - 是否存在一个 c ，使得该 MDP 中的所有策略都是最优策略？举出该例子。
- 如果该 MDP 存在终止状态，这是否会改变问题 a) 的结果？如果是，请给出一个例子并说明理由。

题目 2：REINFORCE 算法优化

为了计算 REINFORCE 的目标估计，需要计算 $(G_t)_{t=1}^T$ 的值，其中 G_t 的表达式是：

$$G_t = \sum_{t'=t}^T \gamma^{t'-t} r_{t'}$$

注意到，计算所有这些值需要 $O(T^2)$ 时间。请给出一种在 $O(T)$ 时间内计算的方法。

题目 3：封冻湖小游戏

在这道题中，需要使用 Gymnasium（前身为 OpenAI Gym）环境实现值迭代和策略迭代。请参照本文件中的示例代码实现。

a) 环境配置

- 使用 Python 3.8 或更高版本完成。
- 完成任务所需依赖包在 `requirements.txt` 中，可以通过 `pip install -r requirements.txt` 安装。
- 建议为创建专用虚拟环境（`virtualenv` 或 `Anaconda`），以避免软件依赖包冲突。
- 编程时，若有任何 `DeprecationWarning` 出现，请忽略。

b) 题目要求

- 编程题 1：阅读示例代码 `algorithm1.py`，并实现 `policy_evaluation`，`policy_improvement` 和 `policy_iteration` 三个函数，迭代停止条件 $\max_s |V_{old}(s) - V_{new}(s)| < 10^{-3}$ ，折扣因子 $\gamma = 0.9$ 。返回最优价值函数和最优策略。
- 编程题 2：实现 `algorithm1.py` 中的 `value_evaluation`，迭代停止条件和折扣因子与编程题 1 相同。返回最优价值函数和最优策略。
- （问题）分别在附件代码中 `Deterministic-4x4-FrozenLake-v0` 和 `Stochastic-4x4-FrozenLake-v0` 中运行上面两小问的方法。在第二个环境中，环境互动是随机的。请从定性和定量两个角度指出随机性是如何影响迭代的次数，以及最终策略。

提示：对于 `Deterministic-4x4-FrozenLake-v0` 环境，前三个状态的值是 0.59, 0.656, 0.729。

作为验证，实现值迭代和策略迭代函数产生的值与它们的误差应该不超过 10^{-3} 。

c) 附件

requirements.txt

```
gymnasium[toy_text]==0.27.0
matplotlib
numpy
scipy
```

algorithm1.py

```
### MDP Value Iteration and Policy Iteration
import argparse
import time

import gymnasium as gym
import numpy as np
from envs import *

np.set_printoptions(precision=3)

parser = argparse.ArgumentParser(
    description="A program to run assignment 1 implementations.",
    formatter_class=argparse.ArgumentDefaultsHelpFormatter,
)

parser.add_argument(
    "--env",
    type=str,
    help="The name of the environment to run your algorithm on.",
    choices=["Deterministic-4x4-FrozenLake-v0", "Stochastic-4x4-FrozenLake-v0"],
    default="Deterministic-4x4-FrozenLake-v0",
)

parser.add_argument(
    "--render-mode",
    "-r",
    type=str,
    help="The render mode for the environment. 'human' opens a window to render. 'ansi' does not render anything.",
    choices=["human", "ansi"],
```

```

        default="human",
    )

"""
For policy_evaluation, policy_improvement, policy_iteration and
value_iteration,
the parameters P, nS, nA, gamma are defined as follows:

    P: nested dictionary of a nested lists
        From gym.core.Environment
        For each pair of states in [1, nS] and actions in [1, nA],
P[state][action] is a
            tuple of the form (probability, nextstate, reward,
terminal) where
                - probability: float
                    the probability of transitioning from
"state" to "nextstate" with "action"
                - nextstate: int
                    denotes the state we transition to (in
range [0, nS - 1])
                - reward: int
                    either 0 or 1, the reward for
transitioning from "state" to
                    "nextstate" with "action"
                - terminal: bool
                    True when "nextstate" is a terminal state (hole
or goal), False otherwise
    nS: int
        number of states in the environment
    nA: int
        number of actions in the environment
    gamma: float
        Discount factor. Number in range [0, 1)
"""

def policy_evaluation(P, nS, nA, policy, gamma=0.9, tol=1e-3):
    """Evaluate the value function from a given policy.

    Parameters
    -----
    P, nS, nA, gamma:
        defined at beginning of file
    policy: np.array[nS]

```

```

        The policy to evaluate. Maps states to actions.
    tol: float
        Terminate policy evaluation when
            max |value_function(s) - prev_value_function(s)|
< tol
    Returns
    -----
    value_function: np.ndarray[nS]
        The value function of the given policy, where
value_function[s] is
        the value of state s
    """

    value_function = np.zeros(nS)

    #####
    # YOUR IMPLEMENTATION HERE #
    #####
    return value_function

def policy_improvement(P, nS, nA, value_from_policy, policy,
gamma=0.9):
    """Given the value function from policy improve the policy.

    Parameters
    -----
    P, nS, nA, gamma:
        defined at beginning of file
    value_from_policy: np.ndarray
        The value calculated from the policy
    policy: np.array
        The previous policy.

    Returns
    -----
    new_policy: np.ndarray[nS]
        An array of integers. Each integer is the optimal action
to take
        in that state according to the environment dynamics and
the
        given value function.
    """

```

```

new_policy = np.zeros(nS, dtype="int")

#####
# YOUR IMPLEMENTATION HERE #
#####
return new_policy

def policy_iteration(P, nS, nA, gamma=0.9, tol=1e-3):
    """Runs policy iteration.

    You should call the policy_evaluation() and policy_improvement()
    methods to
    implement this method.

    Parameters
    -----
    P, nS, nA, gamma:
        defined at beginning of file
    tol: float
        tol parameter used in policy_evaluation()
    Returns:
    -----
    value_function: np.ndarray[nS]
    policy: np.ndarray[nS]
    """

    value_function = np.zeros(nS)
    policy = np.zeros(nS, dtype=int)

    #####
    # YOUR IMPLEMENTATION HERE #
    #####
    return value_function, policy

def value_iteration(P, nS, nA, gamma=0.9, tol=1e-3):
    """
    Learn value function and policy by using value iteration method
    for a given
    gamma and environment.

    Parameters:
    -----

```

```

    P, nS, nA, gamma:
        defined at beginning of file
    tol: float
        Terminate value iteration when
            max |value_function(s) - prev_value_function(s)|
< tol

    Returns:
    -----
    value_function: np.ndarray[nS]
    policy: np.ndarray[nS]
    """

    value_function = np.zeros(nS)
    policy = np.zeros(nS, dtype=int)
    #####
    # YOUR IMPLEMENTATION HERE #
    #####
    return value_function, policy

def render_single(env, policy, max_steps=100):
    """
    This function does not need to be modified
    Renders policy once on environment. Watch your agent play!

    Parameters
    -----
    env: gym.core.Environment
        Environment to play on. Must have nS, nA, and P as
        attributes.
    Policy: np.array of shape [env.nS]
        The action to take at a given state
    """

    episode_reward = 0
    ob, _ = env.reset()
    for t in range(max_steps):
        env.render()
        time.sleep(0.25)
        a = policy[ob]
        ob, rew, done, _, _ = env.step(a)
        episode_reward += rew
        if done:
            break

```

```

env.render()
if not done:
    print(
        "The agent didn't reach a terminal state in {}
steps.".format(
            max_steps
        )
    )
else:
    print("Episode reward: %f" % episode_reward)

# Edit below to run policy and value iteration on different
environments and
# visualize the resulting policies in action!
# You may change the parameters in the functions below
if __name__ == "__main__":
    # read in script argument
    args = parser.parse_args()

    # Make gym environment
    env = gym.make(args.env, render_mode=args.render_mode)

    env.nS = env.nrow * env.ncol
    env.nA = 4

    print("\n" + "-" * 25 + "\nBeginning Policy Iteration\n" + "-" *
25)

    V_pi, p_pi = policy_iteration(env.P, env.nS, env.nA, gamma=0.9,
tol=1e-3)
    render_single(env, p_pi, 100)

    print("\n" + "-" * 25 + "\nBeginning Value Iteration\n" + "-" * 25)

    V_vi, p_vi = value_iteration(env.P, env.nS, env.nA, gamma=0.9,
tol=1e-3)
    render_single(env, p_vi, 100)

```


envs.py

```
import gymnasium as gym
from gymnasium.envs.registration import register

env_dict = gym.envs.registration.registry.copy()
for env in env_dict:
    if "Deterministic-4x4-FrozenLake-v0" in env:
        del gym.envs.registration.registry[env]
    elif "Stochastic-4x4-FrozenLake-v0" in env:
        del gym.envs.registration.registry[env]

register(
    id="Deterministic-4x4-FrozenLake-v0",
    entry_point="gymnasium.envs.toy_text.frozen_lake:FrozenLakeEnv",
    kwargs={"map_name": "4x4", "is_slippery": False},
)

register(
    id="Stochastic-4x4-FrozenLake-v0",
    entry_point="gymnasium.envs.toy_text.frozen_lake:FrozenLakeEnv",
    kwargs={"map_name": "4x4", "is_slippery": True},
)
```